

Machine Learning Assignment 2

Conor Keaney (4BCT)

For this assignment I have chosen to use the CART algorithm to implement and evaluate.

CART

CART (Classification and Regression Tree) algorithm is a Binary decision tree, in that each node can have 0-2 child nodes. Each node represents a value and also a split point on that variable in order to guide values down the tree to the leaf nodes. The leaf nodes contain an output which is used to make a prediction (e.g a target label for classification).

The binary decision tree is created by recursively partitioning the input data in the following steps.

- 1) For all attributes that have yet to be used in the tree, calculate their **Gini impurity** and **information gain** values using the training data.
- 2) Select the attribute with the highest information gain.
- 3) Make a tree node containing that attribute.

This node then will partition the data and this algorithm is applied recursively to each partition.

I have chosen to use **Information Gain** as our metric for this project due to my familiarity with it and it is the primary metric used alongside CART. In order to obtain the Information Gain, normally entropy is used, which is the measure of uncertainty, however as is commonplace when using CART, we will be using the **Gini Index**.

The **Gini Index/Impurity** is calculated by subtracting the sum of the squared probabilities of each class from one. This calculates the amount of uncertainty.

$$\text{Gini} = 1 - \sum_{i=1}^c (p_i)^2$$

Fig.1 Gini Index Formula

Information Gain lets us quantify how much a query reduces that uncertainty given by the Gini Index. It is calculated by taking the impurity of the starting node from the impurity of the two child nodes.

$$\text{Gain} = \text{CurrentGini} - p * \text{Gini(Left)} - (1-p) * \text{Gini(Right)}$$

$$\text{Where } p \text{ is equal to } - \frac{\text{Length of Left branch}}{\text{Length of both left and right branches}}$$

Fig.2 Information Gain

Design decisions

A brief overview of the operation of my program is as follows

Preprocessing

- I used *ctypes* in order to print out a message box to ask the user to add the testing/training files needed for the primary algorithm implementation and the reference implementation.
- *Tkinter* opens the file dialog and the user selects the file(s) required.
- *Ntpath* is used in order to retrieve the name of the file from the filepath, which is then assigned as the filename
- I use my *load_data* method in order to load in the files data into the dataset and also to return the number of entries and to assign and return the list of attributes (i.e. *calorific_value*, *nitrogen*, *turbidity*, *style*, *alcohol*, *sugars*, *bitterness*, *beer_id*, *colour*, *degree_of_fermentation*)
- The data is then shuffled and divided into $\frac{2}{3}$ s training data and $\frac{1}{3}$ testing data using the *split_shuffle* method.

Testing

- Rows of data are then inputted into the *rec_tree* method which is used to recursively go through the tree in order to build the tree by using the *split* function.
- The *split* function is used in order to find the best split for data among all attributes by iterating over all attributes and then calculating the information gain for each in the *gain* function.
- The Information gain (*gain*) is calculated by subtracting the uncertainty starting node with the Gini Index value of two child nodes.
- The Gini impurity value is calculated by implementing the algorithm discussed in Figure 1.
- The *fork* and *compare* functions are used in order to test if the current row passed in by split function is greater than or less than the test value which is used to fork the data, directing down the true branch, or the false branch.
- The *confidence* method is used in order to determine the percentage confidence of the decisions.
- Once the back-end of the tree has been built by these methods, we can now print the tree with the *build_tree* and *print_leaf* methods. The tree is drawn within the console with indentation and arrows in order to mimic a tree structure. At each step it will ask a question and take a step in (e.g. Is degree_of_fermentation > 9.48 ?) with another question for the true branch and a different question for the false branch, which is determined in these earlier functions as discussed above based on their Information gain value.
- In my main function I run a while loop 10 times in order to have 10 different interactions all printed out and from these the average accuracy is taken for both my implementation of the algorithm along with the reference implementation. In this case I used scikit-learn's **DecisionTreeClassifier**.
- The printed output of this was written to a file. This included all 10 trees and percentage accuracies along with the overall average percentage accuracy of the implementation and the reference algorithm.

Results

The outputted results of 100 iterations (Figure 4) of my program and a sample tree (Figure 3) are as follows (10 runs, with 10 iterations per run):

```
Is nitrogen > 0.360463403 ?
--> True:
  Is sugars > 4.015384615 ?
  --> True:
    Predict {'ale': 31}
  --> False:
    Is calorific_value > 44.59734513 ?
    --> True:
      Predict {'stout': 1}
    --> False:
      Predict {'lager': 1}
--> False:
  Is turbidity > 1.800909091 ?
  --> True:
    Is beer_id > 10.31873684 ?
    --> True:
      Is calorific_value > 45.70353982 ?
      --> True:
        Predict {'lager': 1}
      --> False:
        Predict {'ale': 1}
    --> False:
      Predict {'stout': 24}
--> False:
  Is sugars > 3.941538462 ?
  --> True:
    Is degree_of_fermentation > 11.52 ?
    --> True:
      Is bitterness > 17.05 ?
      --> True:
        Predict {'ale': 4}
      --> False:
        Predict {'lager': 4}
    --> False:
      Predict {'stout': 6}
--> False:
  Is bitterness > 19.72 ?
  --> True:
    Predict {'stout': 1}
  --> False:
    Predict {'lager': 29}
```

Fig.3 Sample tree output

Run # (set of 10)	% Accuracy of my implementation	% Accuracy of reference algorithm
1	83.333	80.333
2	84.313	85
3	82.157	82.667
4	85.49	76.667
5	83.333	80
6	82.941	75.333
7	87.255	72.333
8	83.333	80
9	84.117	71.667
10	86.078	74

Fig.4 Accuracy results for 100 iterations

Average % Accuracy of my implementation	Average % Accuracy of reference algorithm
84.235%	77.8%

These details will be outputted to a file called output.txt. In my zip file I will include a sample output file named as such (*sample-output.txt*).

Results can be obtained manually also by running the program and selecting the training / testing files you require..

I used my implementation of DecisionTreeClassifier from Assignment 1 for the reference algorithm. In order to keep correct formatting of the original reference implementation I split data from the beer file into training and testing data files in order to be read in for the x,y testing and training data sets.

Conclusion

This is a simple implementation of the C4.5 algorithm. While I found it very difficult at the beginning however it began to make more sense over time and now I think it was a great learning experience and hopefully one I will carry through to my Final Year Project and in the future workplace.

We can see from the results that the average accuracy of my implementation of the algorithm is higher than my reference implementation of the algorithm.

```

1  import csv
2  import random
3  import sys
4  import pandas as pd
5  from sklearn.tree import DecisionTreeClassifier
6  from sklearn import metrics
7  from tkinter import filedialog
8  import tkinter as tk
9  import ntpath
10 import ctypes
11
12
13 # Used to load in the file data into the dataset
14 # Also returns the number of entries in the file
15 def load_data(file):
16     d = []
17     att = []
18     with open(file, 'r') as f:
19         r = csv.reader(f, delimiter='\t')
20         i = 0
21         for row in r:
22             if i == 0:
23                 att = row
24                 i += 1
25             else:
26                 d.append(row)
27                 i += 1
28     return d, i, att
29
30
31 # Used to shuffle and split the data into 2/3s training data and 1/3 testing data
32
33 def split_shuffle(ds, parts):
34     random.shuffle(ds)
35     p = (len(ds)) // parts
36     test = ds[:p]
37     train = ds[p:]
38     return test, train
39
40
41 # Used to returns all values for a set column
42
43 def get_column(rows, col):
44     cols = []
45
46     for row in rows:
47         cols.append(row[col])
48     return cols
49
50
51 # Used to count the number of each label/feature (beer_style) in the rows passed to the
52 # function
53 def y_count(r):
54     y_num = {}
55     for row in r:
56         y = row[label_col]
57         if y not in y_num:
58             y_num[y] = 0
59         y_num[y] += 1
60     return y_num
61
62
63 # Used to compare and test if the current row is greater than or equal to the test value
64 # in order to split up the data
65
66 def compare(r, test_c, test_val):

```

```

67     if r[test_c].isdigit():
68         return r[test_c] == test_val
69
70     elif float(r[test_c]) >= float(test_val):
71         return True
72
73     else:
74         return False
75
76
77 # Splits the data into two lists for the true/false results of the compare test
78 def fork(r, c, test_val):
79     true = []
80     false = []
81
82     for row in r:
83
84         if compare(row, c, test_val):
85             true.append(row)
86         else:
87             false.append(row)
88
89     return true, false
90
91
92 # Used to calculate the Gini Index/Impurity of the rows inputted (of beer style)
93
94 def gini_index(r):
95     stylesNum = y_count(r)
96     impurity = 1
97
98     for style in stylesNum:
99         style_prob = stylesNum[style] / float(len(r))
100         impurity -= style_prob ** 2
101     return impurity
102
103
104 # Used to calculate the Information gain, incorporates the gini index (impurity)
105
106 def gain(left, right, impurity):
107     p = float(len(left)) / (len(left) + len(right))
108     ig = impurity - p * gini_index(left) - (1 - p) * gini_index(right)
109     return ig
110
111
112 # Used to find the best split for data among all attributes
113
114 def split(r):
115     max_ig = 0
116     max_att = 0
117     max_att_val = 0
118
119     # calculates gini for the rows provided
120     curr_gini = gini_index(r)
121     no_att = len(r[0])
122
123     # Goes through the different attributes
124
125     for c in range(no_att):
126
127         # Skip the label column (beer style)
128
129         if c == label_col:
130             continue
131         column_vals = get_column(r, c)
132
133         i = 0

```

```

134         while i < len(column_vals):
135             # value to compare
136             att_val = r[i][c]
137
138             # Use the attribute value to fork the data to true and false streams
139             true, false = fork(r, c, att_val)
140
141             # Calculate the information gain
142             ig = gain(true, false, curr_gini)
143
144             # If this gain is the highest found then mark this as the best choice
145             if ig > max_ig:
146                 max_ig = ig
147                 max_att = c
148                 max_att_val = r[i][c]
149             i += 1
150
151         return max_ig, max_att, max_att_val
152
153
154     # Used to recursively go through the tree in order to find the optimal attribute to
    split the tree with
155
156     def rec_tree(r):
157         ig, att, curr_att_val = split(r)
158
159         if ig == 0:
160             return Leaf(r)
161
162         true_rows, false_rows = fork(r, att, curr_att_val)
163
164         true_branch = rec_tree(true_rows)
165         false_branch = rec_tree(false_rows)
166
167         return Node(att, curr_att_val, true_branch, false_branch)
168
169
170     # Defines the classifications of the leaf
171
172     class Leaf:
173         def __init__(self, rows):
174             self.predictions = y_count(rows)
175
176
177     # Defines a split node - contains the primary attribute its value and the two child
    branches
178
179     class Node:
180         def __init__(self, att, att_value, true_branch, false_branch):
181             self.att = att
182             self.att_value = att_value
183             self.true_branch = true_branch
184             self.false_branch = false_branch
185
186
187     # Confidence is used in order to determine what is each value
188
189     def confidence(r, node):
190         if isinstance(node, Leaf):
191             return node.predictions
192
193         c = node.att
194         att_value = node.att_value
195
196         if compare(r, c, att_value):
197             return confidence(r, node.true_branch)
198         else:

```



```

261         'beer_id', 'colour', 'degree_of_fermentation']
262
263 ctypes.windll.user32.MessageBoxW(0, "Select your reference algorithm training
data", "File Selection", 0)
264 ref_train_path = filedialog.askopenfilename()
265 train_path_filename = path_name(ref_train_path)
266
267 ctypes.windll.user32.MessageBoxW(0, "Select your reference algorithm testing data",
"File Selection", 0)
268 ref_test_path = filedialog.askopenfilename()
269 test_path_filename = path_name(ref_test_path)
270
271 trainingData = pd.read_csv("training.txt", sep='\t', names=ref_attributes)
272 testData = pd.read_csv("test.txt", sep='\t', names=ref_attributes)
273 sys.stdout = open('output.txt', 'wt')
274
275 # -----#
276 # Main random divisions of the algorithm. Each time the testing and training data is
277 # shuffled and split randomly
278
279 while i < 10:
280     testing, training = split_shuffle(data, 3)
281     tree = rec_tree(training)
282     build_tree(tree)
283
284     correct = 0
285     incorrect = 0
286     for r in testing:
287         print("Actual: %s. Predicted: %s" % (r[label_col], print_leaf(confidence(r,
tree))))
288         for key, value in confidence(r, tree).items():
289             if r[label_col] == key:
290                 correct += 1
291             else:
292                 incorrect += 1
293         print('Percentage Correctly Classified')
294         print(correct / (correct + incorrect) * 100)
295         print('Percentage Incorrectly Classified')
296         print(incorrect / (correct + incorrect) * 100)
297
298     i += 1
299     avg_acc += correct / (correct + incorrect)
300
301 # -----#
302 # REFERENCE IMPLEMENTATION
303
304 x_train = trainingData[featuredCols]
305 y_train = trainingData.beer_style
306 x_test = testData[featuredCols]
307 y_test = testData.beer_style
308
309 dtc = DecisionTreeClassifier()
310 dtc = dtc.fit(x_train, y_train)
311 y_predict = dtc.predict(x_test)
312
313 avg_ref_acc += metrics.accuracy_score(y_test, y_predict)
314 print('Reference Algorithms Percentage Accuracy')
315 print(metrics.accuracy_score(y_test, y_predict)*100)
316 # -----#
317
318 print("\nThe Average Accuracy across 10 iterations: ")
319 acc10 = avg_acc / 10 * 100
320 print(acc10)
321
322 refAcc = ((avg_ref_acc / 10) * 100)
323 print("\nThe Average Accuracy for the reference decision tree classifier across 10
iterations: ")

```

```
324     print(refAcc)
325
```